

# LLMs for Preparing Data in R

*Guest lecture for ETC5512: Wild Caught Data*

Cynthia A. Huang

19 May 2025




# Introduction

# Learning goals

- develop your understanding of:
  - what LLMs are,
  - different types of “wild caught data” tasks that LLMs can help with,
  - how to use and check LLMs for specific data preparation tasks,
  - how to interact with LLMs in R using {ellmer}

# Lecture outline



- About Me!
- About LLMs
- Using LLMs for “Wild Caught Data” tasks
- Using LLMs in R with `{ellmer}`



**ANTHROPIC**

**About Me!**

# Who am I?

- 🎓 Research Fellow in EBS supervised by *Prof. Rob Hyndman*
- 🎙️ Regular host on [The Random Sample](#) podcast
- 🧑 Ex-Tutor for Wild Caught Data (2020)
- 💵 Previously:
  - Economics at the University of Melbourne
  - Catching wild data for empirical researchers:
    - wikipedia entries, archival magazines, trade databases, satellite images, online retail prices...

# What do I work on?

- 🎓 Thesis: **Unified Statistical Principles and Computational Tools for Data Harmonisation and Provenance**, supervised by:
  - *Prof. Rob Hyndman* (EBS, MBUS), *Prof. Simon Angus* (Econ, MBUS), and *Dr. Sarah Goodwin* (HCC, FIT)
- 📊 Research Interests
  - 🍌 Designing tools and workflows for wild caught data!
  - 🤖 Leveraging LLMs and genAI for data wrangling and cleaning
    - 📎 using LLMs to correct manual data entry errors – Current MBAT research internship!

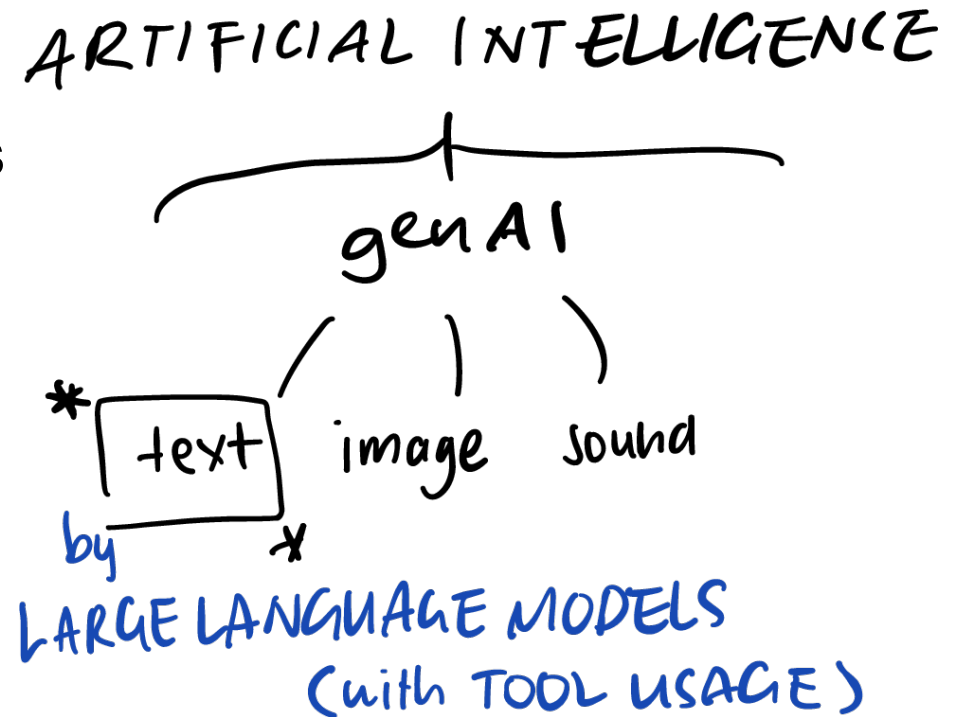
# About Large Language Models



# Generative AI and LLMs

*Generative AI* refers to:

- computer algorithms and systems
- that can generate content such as text, images and sound
- based on patterns learnt from existing data




# What are Large Language Models?

LLMs are...

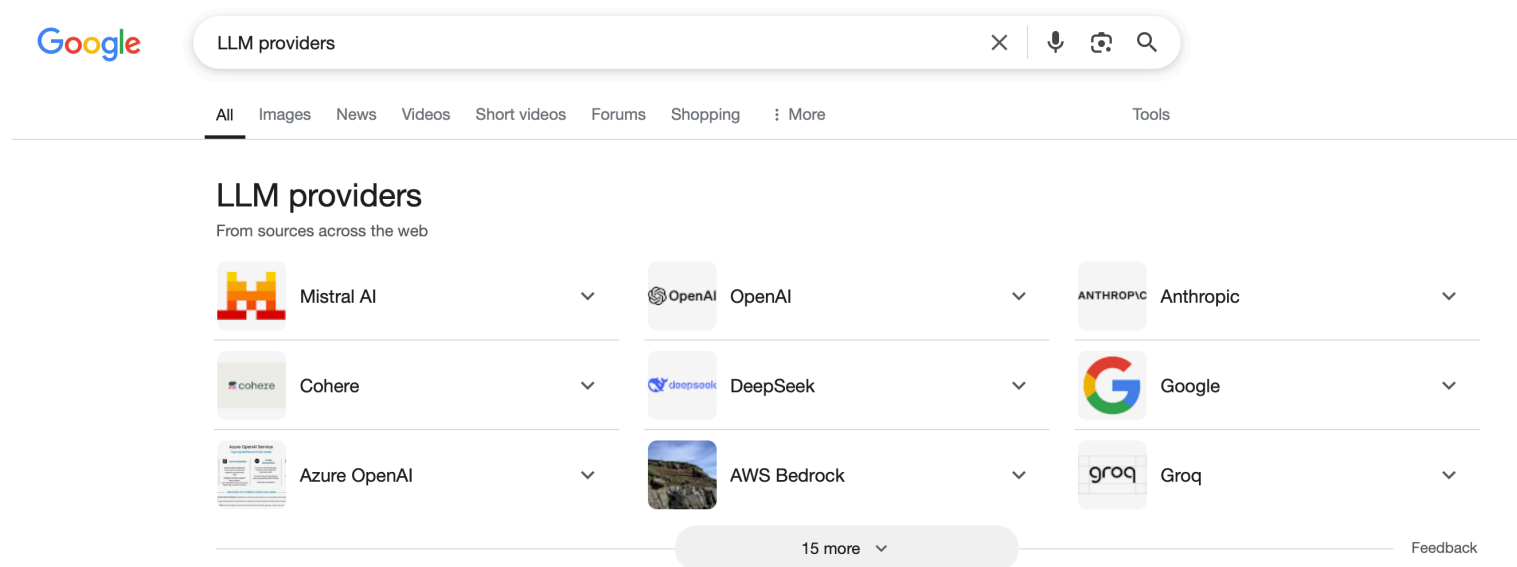
- code writers?
- encyclopedias?
- assignment help?
- translators?

We often understand tools by what they can do for us, not how they work.



# Who makes LLMs?

**LLM providers** develop and offer access to large language models and systems



Today, we will see demos of *Anthropic* and *OpenAI* models.

# Why are there so many different models?

LLM providers offer paid and free access to multiple models:

- **OpenAI:** GPT-3 and 4, o-series,
- **Anthropic:** Claude Haiku, Sonnet and Opus
- **Google:** Gemini Flash and Pro
- **Meta:** Llama 3, Llama 4
- **Alibaba:** Gwen 2.5, 3, Max, Plus and Turbo

Different models are designed to be good at different things:


- chain-of-thought reasoning vs. instruction following
- multimodal support: images, audio, video AND text
- multilingual processing: translation, content generation
- specific domains: medicine, finance, legal

# Model differentiation

Learn more about picking the right tool:

- 🗣️ [Beyond ChatGPT: THE RAPIDLY EVOLVING LANDSCAPE OF AI](#)
- 💡 [Suggestions on provider/model choice in 'ellmer' docs](#)
- 📖 [OpenAI Model Selection Guide](#)
- 📖 [Anthropic Claude Model Comparison Table](#)

# How can we interact with LLMs?



**OpenAI**

**ANTHROPIC**


## 1. Web-based Chat Interface

- ChatGPT, Claude AI, Qwen Chat
- offers additional formatting of outputs, access to tools, other 'quality of life' features

## 2. Programmatic Interfaces

- requires an API key
- interact using code with an LLM from within an R session

# How can we interact with LLMs?




**OpenAI**

**ANTHROPIC**

- 1. Web-based Chat Interface**
- 2. Programmatic Access**
- 3. Other interfaces**

- mobile (chat) applications
- voice assistants
- embedded LLMs (e.g. suggestions in Gmail)

# How can we interact with LLMs?



**OpenAI**

**ANTHROPIC**

1. Web-based Chat Interface
2. Programmatic Access
3. Other interfaces

Today, we will use web-based chat via:

- [chatgpt.com](https://chatgpt.com)
- [claude.ai](https://claude.ai),

as well as programmatic access via:

- `{ellmer}`



# How do we verify what LLMs are doing?

For a dishwasher, we consider:

- ...are the dishes clean?
- is there any dirt on the dishes?


For LLMs...?

It depends on the (data) task!



# **Using LLMs for “Wild Caught Data” Tasks**

# What are common WCD tasks?



**i Discuss in groups [5 mins]**

- What tasks are involved in preparing Wild Caught Data for analysis? *List at least 3.*
- Which of these might be more or less suitable for addressing with LLMs? Why?

# Using LLMs for Data Preparation Tasks

## 1. Generating data wrangling code

- 'indirect' use of LLMs to clean data

## 2. Transforming and generating new data

- creating example datasets (e.g. when documenting a custom R function)
- changing data between formats (csv to json)

# **Using LLMs for Data Preparation Tasks**

**1. Generating data wrangling code**

**2. Transforming and generating new data**

**3. Modifying and augmenting existing data**

- filling in missing data
- correct typos or inconsistencies
- creating new columns based on existing ones

# Modifying and augmenting data

- filling in missing data
  - **'look up' facts:** author nationality: `Jane Austen`
  - **suggesting values:** missing volume units for drinks: `300?`
- correct typos or inconsistencies
  - **harmonise different abbreviations:** `{Victoria, VIC, Vic} → {VIC}`
- creating new columns based on existing ones
  - **comparison and categorisation:** are `teachers` and `instructors` similar occupations?
  - **summarise text:** key points in free-form survey responses
  - **extract info:** name of the movie in a film review

# Prompts for Data Preparation Tasks

- prompts need to include:
  - instruction and data!
- responses would ideally:
  - return data of the expected type
  - and in a easy to import format

## Try yourself! [10mins]

1. Pick a starting prompt from the next page
2. Fill in the necessary data.
3. Try the prompt in your choice of LLM chat (e.g. ChatGPT, Qwen Chat, Claude AI etc.).
  - What output did you get in return?
  - Could you import it into R easily?
4. Modify the prompt to return the answer in a more useful format.

# Starting Prompts

- **‘look up’ facts:** “What nationality is *the author* `<author name>`?”
- **suggesting values:** “What is the likely volume unit *of a beverage* `<can>` *of with* a volume of `<300>`?”
- **harmonise different abbreviations:** “Convert the following list of Australian states to all use three-letter state codes: `<list>`”
- **comparison:** “How similar are these two occupations: `<occupation A>`, `<occupation B>`?”
- **summary:** “Summarise the following survey response: `<text>`”
- **extraction:** “What movie is the follow review about”: `<review text>`



# Requesting different output formats

- LLM can respond in many different ‘text’ formats.
- Some are more useful than others.

Let’s look at an [example conversation with ChatGPT](#) for the following prompt

Convert the following list of Australian states to all use three-letter state codes (e.g. VIC, TAS):

- Victoria
- NSW
- N.T.
- ACT
- Queensland

# Verifying success

**Verification** is the most important skill when using LLMs, but it requires:

- Clearly defined tasks and expected outcomes
- Ways of checking the outcomes have been achieved






# Approaches to verification

There are multiple ways to verify outcomes match expectations.

- **Positive verification:** Define characteristics of 'success'
- **Negative verification:** Figure out signals or signs of 'failure'
- **Trust-based verification:** Seek assurance and confirmation of 'success'






# Checking on the dishes

There are multiple ways to verify outcomes match expectations.

- **Positive verification:** Are the dishes clean?
- **Negative verification:** Is there any dirt on the dishes?
- **Trust-based verification:** Ask the machine if the dishes are clean...?





# Using LLMs in R with {ellmer}



# Beyond web-based interfaces

Different interfaces mean different data preparation workflows:

- LLM web-interface = copy/paste
- programmatic interfaces = code and variables

Using the {ellmer} R package we can:

- send prompts to that LLM from an R session
- construct prompts from with imported data
- systematically docutest different prompts BEFORE scaling up
- manipulate response content using code

# Connecting {ellmer} to an LLM

The basic steps:

- Installing `{ellmer}`
- Getting an API key from the LLM provider you want to use > knock! knock! who's there? Cynthia's R session!
- Storing the API key where ellmer can find it
- Starting a chat session using the relevant `ellmer::chat_*`()

Today, we'll interact with the OpenAI API via `ellmer::chat_openai()`

# LIVE DEMO: Chatting via {ellmer}

```
1 library(ellmer)
2 ## A session is like a chat conversation
3 session <- chat_openai()
4
5 question <- "
6   How can I pick a random letter from A-Z.
7 "
8
9 ## send a question to the 'chat'
10 session$chat(question)
11
12 ## clarify your request
13 session$chat("Return R code only")
14
15 ## inspect all turns in the session so far
16 session
```

What if we *always* want the LLM to return R code?

# LIVE DEMO: System Prompts

```
1 library(ellmer)
2
3 session_tidy_expert <- chat_openai(system_prompt = "
4   You are an expert R programmer
5   who prefers the tidyverse.
6   Only return code without explanation.
7 ")
8
9 session_tidy_expert$chat(question)
10
11 session_tidy_expert
```

*Example adapted from [ellmer docs](#)*

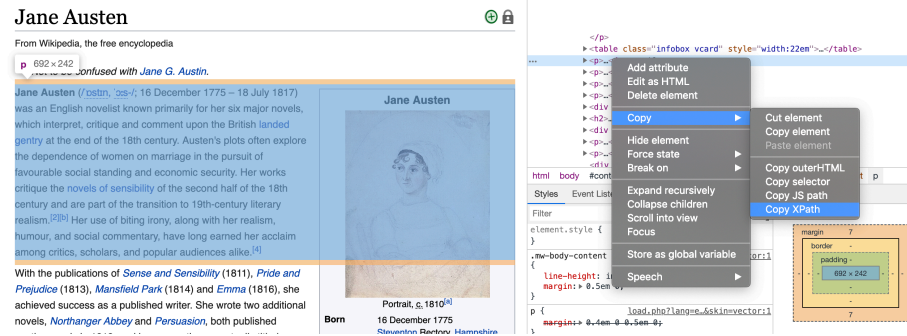
# Sessions and system prompts

- A **chat session** is a single conversation instance between a user and an LLM
- A **R session** is an active workspace where you're running the R programming language
- A **system prompt** is the behind-the-scenes instruction manual that tells an AI assistant:
  - what tone to use,
  - what information the system can access,
  - and how to handle different types of questions or requests

# Revisiting Author Nationalities

Could we use an LLM to extract Jane Austen's nationality?

How can an LLM help with missing nationalities?



WCD 2025, S1, Week 10 - Slide 29

```
# A tibble: 11 × 3
  author_name      nationality
author_links      <chr>
<chr>
1 Rick Warren      <NA>
https://en.wikipedia.org/wiki/R... <NA>
2 William Griffith Wilson <NA>
https://en.wikipedia.org/wiki/B... <NA>
3 F. Scott Fitzgerald <NA>
https://en.wikipedia.org/wiki/F... <NA>
4 John Green        <NA>
https://en.wikipedia.org/wiki/J... <NA>
5 Sam McBratney     <NA>
https://en.wikipedia.org/wiki/S...
```

# LIVE DEMO: Extract Nationalities

```
1 text <- "Jane Austen (/ˈɒstɪn, ˈɔːstɪn/ OST-in, AW-stin; 16 December 1775 – 18 July 1817) w
2
3 session_read <- chat_openai(system_prompt = "You are a data entry assistant.")
4
5 session_read <- chat_openai()
6
7 nationality_prompt <- "Nationality of person"
8
9 session_read$extract_data(text, type = type_string(description = nationality_prompt))
10
11 std_prompt <- "Extract structured data of the nationality of person. Return only ISO 3-digi
12
13 session_read$extract_data(text, type = type_string(description = std_prompt))
```

What if we don't have the extended text description available?

# LIVE DEMO: Ask for Nationalities

```
1 library(dplyr)
2
3 author_df <- readr::read_csv('example_data/week10-author_df.csv')
4
5 short_prompt <- "Nationality of person only"
6 session_lib <- chat_openai(system_prompt = "You are a librarian with expert knowledge of po
7
8 ## let's ask about multiple authors
9 author_df |>
10   tail(6) |>
11   rowwise() |>
12   mutate(nationality_llm =
13     session_lib$clone()$extract_data(author_name,
14                                     type = type_string(short_prompt))
15   )
```



# Evaluation through agreement

Another way to verify data quality is via **consensus**.

Here are nationalities returned by OpenAI's gpt-4o model on 18/05/2025:

```
# A tibble: 6 × 3
  author_name      nationality nationality_llm
  <chr>            <chr>          <chr>
1 Agatha Christie English        British
2 J. R. R. Tolkien English        British
3 Vladimir Nabokov Russian        Russian-American
4 Umberto Eco      English        Italian
5 Shere Hite       German        American
6 J. P. Donleavy   Irish         American-Irish
```

How could you use this information to assess your data quality?

# How much do LLMs 'know'?

What happens if we ask about less widely-known people?

```
1 session <- chat_openai()  
2 session$chat("List the instructors of Monash University's wild caught data course.")
```

Let's try again via the [ChatGPT web interface](#)

```
"List the instructors of Monash University's wild caught data course."
```

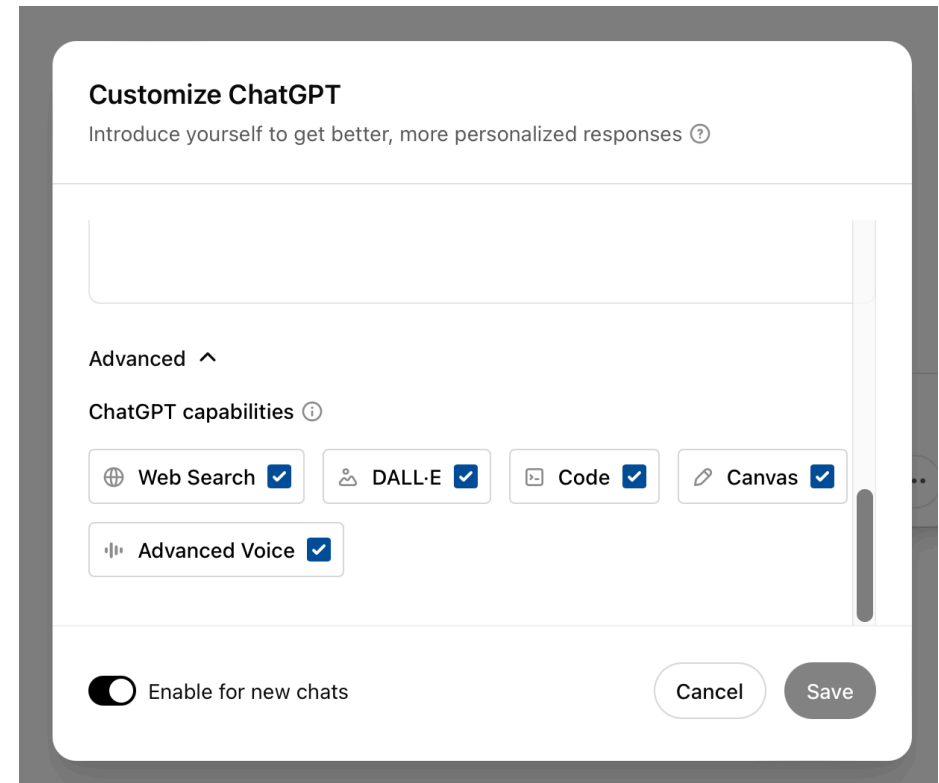
Why do you think ChatGPT able to answer this request via the web interface in this [demo chat](#)?

# LLM agents, web search and more

LLMs can be integrated with additional capabilities such as:

- search
- image generation
- gmail, gdrive and gcal search (e.g. in Claude AI)

How do you think the ability to search affects verifying task success?



# **Final Comments**

# Ethics and AI safety

## Generative AI acknowledgement

I used AI in the following ways:

- i. generate definitions and suggested explanations for key concepts covered in this lecture. I used Claude AI to suggest definitions for terms like 'Generative AI', and 'System Prompt', and to generate lists of "top LLM providers in 2025" and "ways of interacting with LLMs".

# Key takeaways



There are many LLM models and systems which:

- generate text outputs: code and ‘data’
- call tools to get more (text) input
- are available via different types of user interfaces: chat vs. programmatic

When using LLMs for preparing data, think about:

- Breaking your larger, overall data preparation goal into *specific* tasks
- Ways to verify the LLM’s performance on your task

# What we've learnt...



- basics of what LLMs can do, who provides them and what differentiates them
- different 'wild caught data' tasks that LLMs can be used for
- how to approach verifying LLM output for data preparation tasks
- how to interact with LLMs from R using {ellmer}